

ACADEMIC TASK-2

CSE316

(OPERATING SYSTEMS)

COMPUTER SCIENCE AND ENGINEERING

SUBMITTED BY:

NAME: *Shikher singh*

REGISTRATION NO.

ROLL NO: **31**

SECTION:K24PR

SUBMITTED TO:

KUNDAN



LOVELY
PROFESSIONAL
UNIVERSITY

LOVELY PROFESSIONAL UNIVERSITY

INDEX

Sr.No	Content	Page No.
1.	Project Overview	
2.	Module-Wise Breakdown	
3.	Functionalities	
4.	Technology Used 1. Programming languages 2. Libraries and Tools 3. Other Tools(GitHub,etc)	
5.	Flow Diagram	
6.	Revision Tracking on Github	
7.	Conclusion and Future Scope	
8.	References	
9.	Appendix 1. Appendix A: AI Generated project Breakdown Report 2. Appendix B: Problem Statement 3. Appendix C: Solution / Complete Code	

1. Project Overview:

Multiprogramming systems allow multiple processes to run simultaneously, sharing CPU and memory resources. However, static resource allocation can lead to inefficiencies such as CPU starvation, memory overuse, bottlenecks, and performance degradation.

This project implements an Adaptive Resource Allocation System, which dynamically monitors real-time performance metrics and automatically adjusts CPU and memory allocation to optimize system utilization.

Goals:

- Continuously monitor CPU/memory usage of multiple processes.
- Detect overload conditions and prevent system bottlenecks.
- Dynamically reassign resources (CPU shares, memory limits, priorities).
- Improve overall system throughput and responsiveness.
- Provide a dashboard for visualization and manual override.

Expected Outcomes:

- Working monitoring daemon.
- Intelligent allocation engine (heuristic-based, with optional ML).
- Visualization dashboard displaying real-time metrics and decisions.
- Improved system stability and resource utilization.

Scope:

In-scope:

- Real-time monitoring
- Dynamic CPU/memory allocation
- Logging & visualization
- ML-based advisor (optional)

Out-of-scope:

- Kernel-level scheduler modification
- Cluster or distributed system management
- Advanced container orchestration

2. Module Wise Breakdown:

The project is divided into three core modules, each with a well-defined purpose, input-output interaction, and internal components.

These modules work together to implement real-time adaptive resource allocation in a multiprogramming system.

Module A — Monitoring & Data Collection Module:

1. Purpose

This module acts as the system's sensory unit, continuously observing CPU, memory, and process-level statistics. It collects raw data from the operating system at regular intervals and makes it available for the allocation engine and dashboard.

2. Responsibilities

- Real-time CPU usage tracking (system-wide & per process).
- Memory consumption measurement (RSS, virtual memory, swap usage).
- Identification of resource-intensive processes.
- Detection of threshold violations or sudden usage spikes.
- Logging time-series data for trend analysis.
- Exposing a standard data interface for the Allocation Engine.

3. Key Components

A. Metric Collector

Uses psutil (Python) or OS utilities to gather:

- CPU %
- Memory %
- Process table
- Process priority (nice value)
- Load average

B. Threshold Evaluator

Evaluates pre-defined thresholds, such as:

- CPU > 90% for X seconds
- Memory usage > 80%
- Swap usage > 20%

Triggers alerts for the Allocation Engine.

C. Data Exporter

Provides collected data:

- through APIs (Flask/FastAPI),
- through Prometheus metrics,
- or via shared memory/log files.

4. Data Flow

OS → Metric Collector → Evaluator → Data Exporter → Allocation Engine & Dashboard

Module B — Allocation Engine (Policy Manager + Controller)

1. Purpose

This module is the decision-making brain of the system.

It analyzes the collected data, determines whether resource adjustments are required, and applies those changes dynamically.

2. Responsibilities

- Apply scheduling policies and algorithms.
- Analyse system behaviour and prioritize processes.
- Prevent bottlenecks and maintain system stability.
- Adjust CPU share, memory limits, and process priority in real-time.
- Maintain fairness among active programs.
- Implement rollback if the new allocation reduces performance.

3. Key Components

A. Policy Manager (Decision Logic)

Implements the resource allocation strategy using either:

a. Heuristic Rules

Examples:

- If a process consistently exceeds CPU threshold → limit CPU using cgroups.
- If a critical process is starved → increase its CPU share or priority.
- If memory usage is too high → lower the memory limit of background processes.

b. ML-based Advisor (Optional)

ML Models can:

- Predict future workload spikes,
- Recommend optimal allocations,

- Classify processes into categories (critical / background / optional).

B. Performance Analyzer

Evaluates:

- Overall system throughput
- Process waiting time
- CPU idle percentage
- Memory fragmentation or swapping

C. Resource Controller (Actuator)

Applies decisions using OS-level tools:

- cgroups v2 for CPU and memory limits
- nice/renice for priority
- RLIMIT for memory ceiling
- I/O control if required

D. Safety Manager

Ensures:

- Cooldown intervals between adjustments
- Minimum CPU/memory limits
- Rollback mechanism if performance worsens
- Logging of all changes

4. Data Flow

Monitoring Data → Policy Manager → Controller → OS Resource Manager → Updated Allocation

Module C — Dashboard & Visualization Module

1. Purpose

Provides a user-friendly interface for monitoring live system performance, resource allocation decisions, and the status of running processes.

2. Responsibilities

- Display real-time CPU and memory graphs.
- Show per-process statistics in a table.
- Log controller decisions with timestamp and rationale.
- Allow manual override of system allocation.

- Provide downloadable reports for system analysis.

3. Key Components

A. Real-Time Visualization Engine

Visual components created using:

- Chart.js
- Plotly
- Grafana dashboards
- Custom React/Flask frontend

Shows:

- Live CPU Usage Graph
- Memory Utilization Heatmap
- Process Resource Table
- Allocation Events Timeline

B. Backend API

Serves monitoring data and allocation events to the frontend.

Could be developed using:

- Flask
- FastAPI

C. Manual Override Interface

Allows the system operator to:

- Increase/decrease CPU shares manually
- Kill or pause processes
- Manually mark a process as critical or low priority

D. Log Viewer

Shows:

- Allocation decisions
- Threshold violations
- Errors and warnings
- ML predictions (if enabled)

4. Data Flow

Monitoring Module → API Backend → Dashboard UI → User Actions → Allocation Engine

3. Functionalities:

The Adaptive Resource Allocation System includes a rich set of functionalities distributed across three major modules. Each functionality is designed to improve system performance, optimize resource usage, and ensure smooth multiprogramming execution.

Module 1: System Monitoring & Resource Control – Detailed Functionalities

This module forms the core operational engine of the project. It continuously monitors programs and dynamically adjusts CPU and memory resources.

1. Real-Time Process Monitoring

- Tracks all active processes and threads in the system.
- Continuously collects metrics such as:
 - CPU utilization percentage
 - Memory consumption (RSS, VMS)
 - I/O statistics (read/write bytes)
 - Context-switch rate
 - Process priority and scheduling class
- Data is updated at regular intervals (e.g., every 1 second) using tools like *psutil*.

Why Needed:

Helps in identifying high-demand processes, overloaded resources, and overall system load trends.

2. Performance Metrics Logging

- All collected data is logged in time-series format.
- Logs include timestamps, process IDs, usage spikes, and critical events.
- This data is used for:
 - Machine learning training
 - Trend analysis
 - Visual analytics

Example:

Logs can reveal that a process shows increasing memory usage over time → probable memory leak.

3. Dynamic CPU Allocation

- Adjusts CPU shares/weights or time slices for processes.
- Supports multiple strategies:
 - Time-slice rebalancing
 - Dynamic priority boosting
 - Fair-share scheduling
- Ensures that no single process monopolizes CPU time.

Example:

If Process A is interactive (foreground), its CPU quantum is increased automatically.

4. Dynamic Memory Allocation & Reclamation

- Monitors memory pressure and reallocates as required.
- Identifies idle or inactive processes and reclaims unused memory.
- Prevents memory overload by:
 - Limiting runaway memory usage
 - Terminating misbehaving or crashed processes
- Ensures fair distribution according to process importance.

Example:

If a background process consumes excessive memory, it is throttled or assigned a smaller memory limit.

5. Bottleneck and Anomaly Detection

- Detects abnormal conditions such as:
 - Sudden CPU spikes
 - Memory leaks
 - High waiting/blocked time
 - Thrashing (excessive paging)
- When detected, triggers corrective actions:
 - Relocating resources
 - Adjusting priorities
 - Issuing alerts to the UI

Scenario:

If system CPU usage > 95% for too long → pause non-critical tasks.

6. Custom Scheduling Policies

- Implements user-defined or dynamic scheduling algorithms:
 - Round Robin
 - Priority Scheduling
 - ML-based scheduling
 - Reinforcement-learning-driven resource allocation
- Scheduler communicates with ML engine for intelligent decisions.

Benefit:

Provides smarter decisions than static OS scheduling.

Module 2: Predictive Analytics & ML-Based Engine – Detailed Functionalities

This module provides intelligence and future insight into system behavior, making allocation adaptive rather than reactive.

1. Load Prediction

- Predicts future CPU and memory usage using:
 - Linear Regression
 - Random Forest
 - Time-series forecasting
- Helps the system prepare for workload spikes before they occur.

Example:

If a process consistently increases memory every second → ML predicts a memory spike in 30 seconds.

2. Process Behavior Analysis

- Learns from historical performance patterns.
- Distinguishes between:
 - CPU-bound processes
 - I/O-bound processes
 - Memory-heavy tasks
 - Idle programs
- Allocates resources accordingly.

Benefit:

CPU-bound tasks get more CPU; I/O tasks get system priority to reduce wait times.

3. Anomaly/Misbehavior Detection

- Identifies processes showing unwanted behavior:
 - Memory leaks
 - Sudden resource spikes
 - Unresponsive/hanging state
- Uses Isolation Forest or statistical outlier detection.

Action Taken:

Alert is raised in the UI, or the system throttles/kills the offending process.

4. Policy Optimization System

- Produces optimal resource distribution based on algorithmic heuristics or ML.
- Computes:
 - Which process needs more CPU
 - Which process can be downgraded
 - Whether current scheduling is fair
- Generates actionable decisions for Module 1.

Example:

If three processes are competing for CPU, ML allocates:

- High priority → interactive process
- Moderate → batch job
- Low → background sync task

5. Real-Time Decision Engine API

- Exposes ML-generated decisions to other modules.
- Supports:
 - JSON-based API calls
 - Inter-thread communication
 - Event-driven triggers
- Ensures high-speed synchronization between monitoring and UI.

Module 3: User Interface & Visualization Dashboard – Detailed Functionalities

Provides real-time insights and user control over the system.

1. Live Resource Usage Dashboard

- Displays continuous updates of:
 - CPU usage graphs
 - Memory utilization
 - Per-process statistics
 - I/O performance
- Graphs auto-refresh every second.

Tools: Matplotlib / Plotly / PyQtGraph

2. Process Control Panel

- Allows the user to:
 - Increase/decrease process priority
 - Kill or restart a process
 - Freeze or pause a task
 - Allocate CPU/memory manually

Example:

If a user wants a specific program to finish quickly → increase its priority/CPU share.

3. Prediction Visualization

- Shows ML output such as:
 - “Predicted CPU spike in Process 4 in 20 seconds”
 - “Memory leak detected in Process App.exe”
- Helps users understand system behavior beforehand.

4. Alerts and Notifications

- Sends alert pop-ups when:
 - System reaches critical load
 - Bottlenecks are detected
 - ML finds anomalies
- Alerts are color coded (Green/Yellow/Red).

5. System Logs & Reports

- Displays a table/log viewer for:

- Past warnings
- Reallocation events
- CPU/memory usage history
- User actions

Export options: CSV, TXT, PDF

6. Settings & Customization

Users can customize:

- Monitoring interval
- Thresholds for alerts
- Resource allocation limits
- Visualization theme
- ML model settings

7. Real-Time Simulation Environment

If no real processes are available, UI can run a simulated multiprogramming environment, showing:

- Fake processes
 - Simulated loads
 - Adaptive decisions
- Useful for students and demonstration.

4. Technology Used

Programming Languages

- Python
- Optional: C/C++ (for low-level process interaction, OS simulation)

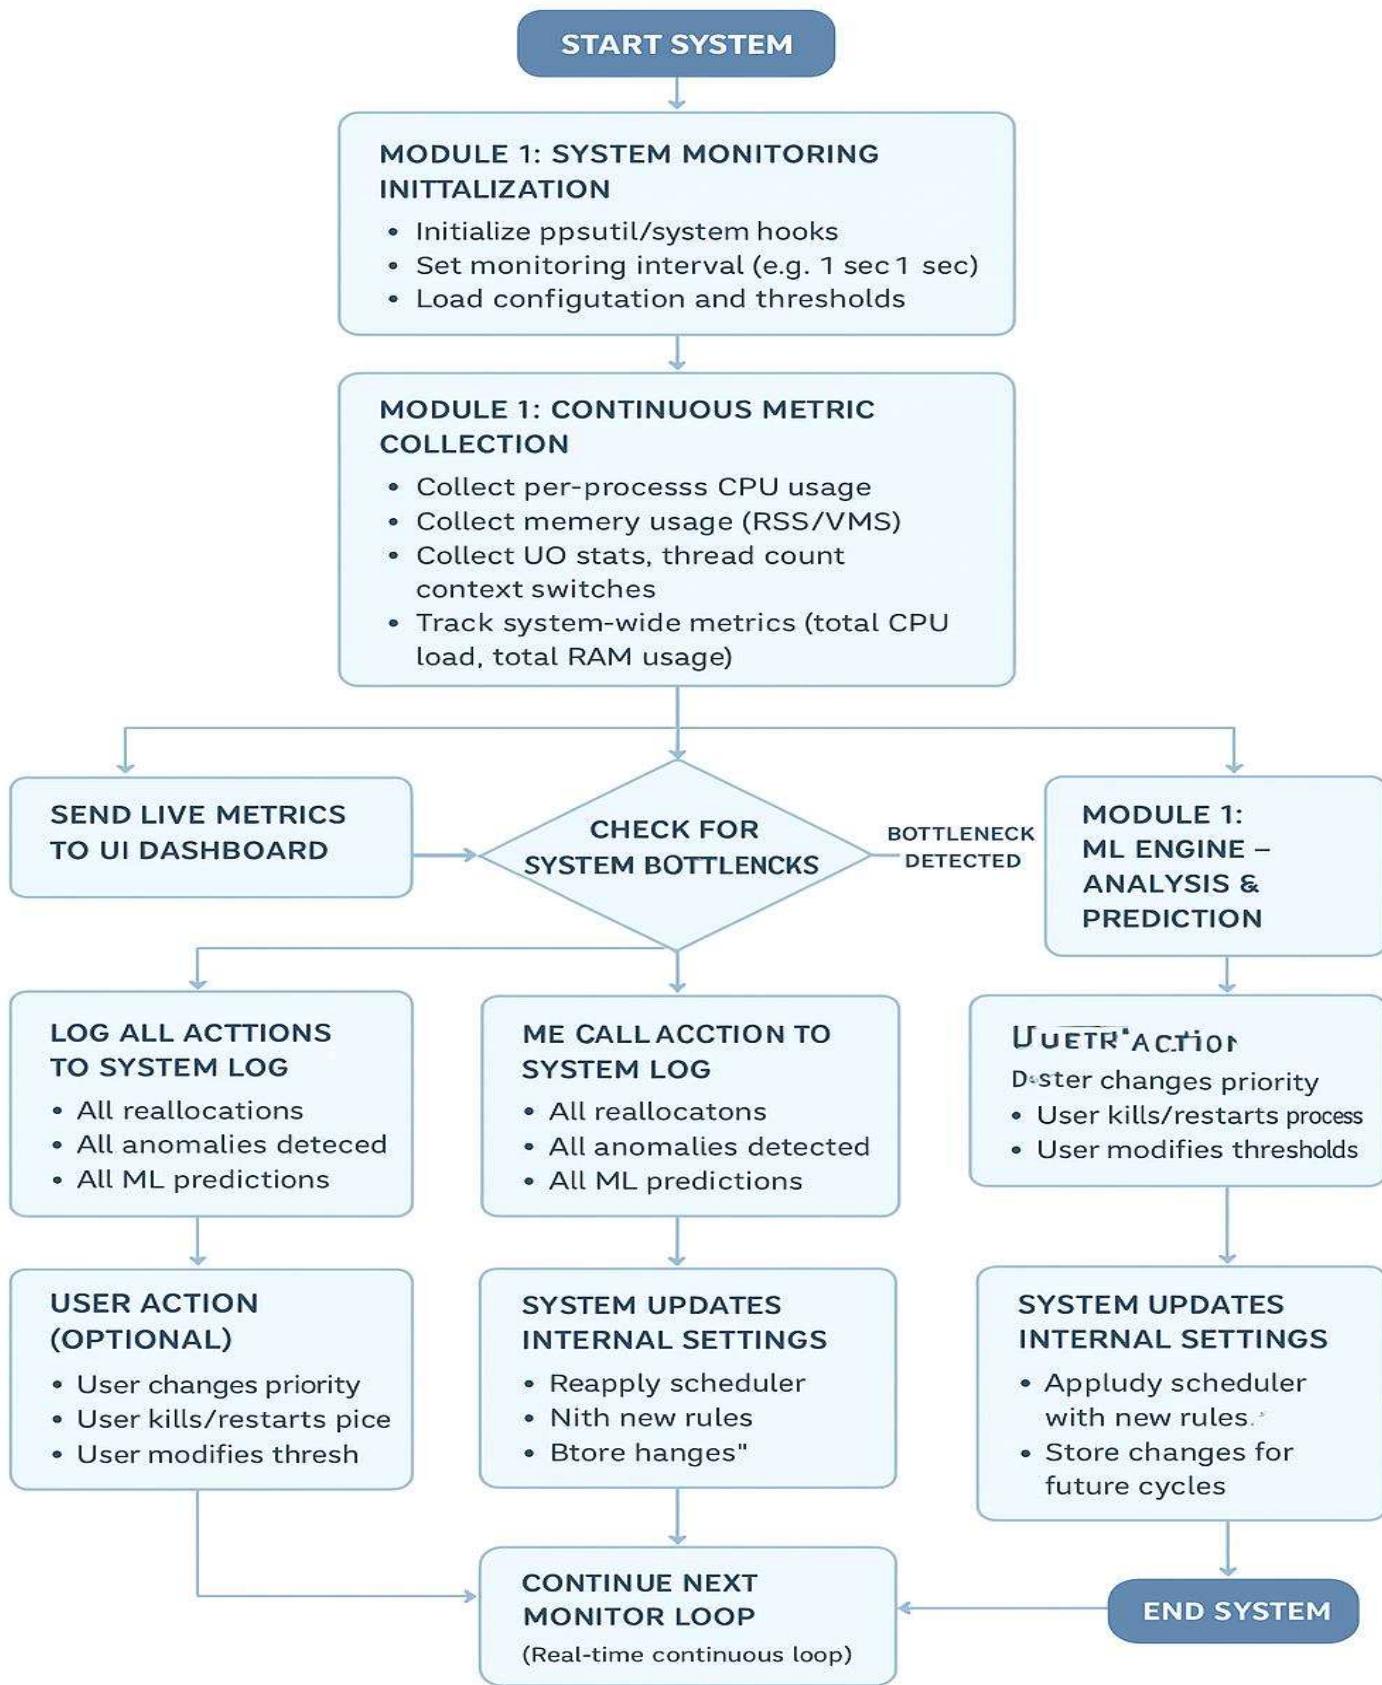
Libraries and Tools

- psutil – system monitoring
- Matplotlib / Plotly – visualization
- scikit-learn – ML models
- pandas, numpy – analytics
- Tkinter/PyQt/Streamlit – UI
- threading / asyncio – concurrency
- logging – event logs

Other Tools

- GitHub – version control
- VS Code / PyCharm – development
- Jupyter Notebook – ML model testing
- Draw.io – flow diagrams

5. Flow Diagram:



6.Revision Tracking on GitHub:

Repository link: <https://github.com/Shikhsingh01/adaptive-resource-allocation>

7. Conclusion and Future Scope:

Conclusion:

The project “Adaptive Resource Allocation in Multiprogramming Systems” successfully demonstrates an intelligent, automated approach to managing system resources such as CPU and memory in dynamic computing environments. Traditional static algorithms often fail to handle fluctuations in workload intensity or unpredictable process behavior. In contrast, the adaptive system designed in this project combines real-time system monitoring, machine-learning-based prediction, and dynamic resource control, providing a smarter and more efficient solution.

Through continuous metric collection, the system identifies performance bottlenecks such as high CPU load, memory saturation, or I/O delays. The ML engine enhances decision-making by predicting future resource requirements and detecting anomalous activities like memory leaks. Based on this intelligence, the system optimally reallocates resources—ensuring fairness, preventing starvation, improving response time, and maximizing overall throughput.

The user interface adds further value by offering intuitive visualization, real-time dashboards, alerts, and manual control options. Together, these features create a complete ecosystem that balances automation with user flexibility. Overall, the project demonstrates that adaptive, intelligent resource allocation significantly improves system stability, efficiency, and performance in multiprogramming environments.

This work lays a strong foundation for more advanced autonomous resource management systems that mirror modern operating system schedulers and cloud orchestration platforms.

Future Scope:

While the system presents a functional and intelligent model for adaptive resource allocation, there is significant potential for enhancements and real-world deployment. Future improvements can include:

1. Integration with Kernel-Level Scheduling

Current implementation operates at the user/application level. Future iterations can:

- Integrate directly with OS kernel modules
- Modify actual CPU scheduler parameters
- Influence memory paging, swapping, and caching decisions

This would make the system suitable for real operating system development.

2. Implementation of Deep Learning Models

Machine learning can be expanded with:

- LSTM/GRU models for advanced time-series forecasting
- Reinforcement Learning (RL) for real-time policy optimization
- Neural networks for behavior classification

Such models would further enhance prediction accuracy and adaptability.

3. Support for Distributed Systems

Modern applications run across clusters, VMs, and containers. Future versions can:

- Extend resource monitoring to distributed nodes
- Implement load balancing between multiple servers
- Create a cloud-based adaptive resource allocator

This aligns with cloud computing and microservices architecture.

4. Container and Virtual Machine Resource Management

The system can be adapted to manage containerized environments by integrating with:

- Docker API
- Kubernetes (K8s) scheduler
- Hypervisors like VMware, KVM, or VirtualBox

This would make it highly relevant for DevOps and cloud-native platforms.

5. Energy-Aware Resource Allocation

Future models can consider power consumption, making the system suitable for:

- Mobile devices
- Green computing
- Data center energy optimization

AI-driven decisions could minimize energy usage while maintaining performance.

6. Fault-Tolerant and Self-Healing Capabilities

The system can be enhanced to:

- Recover automatically from process failures
- Detect and mitigate system crashes
- Restart failed modules or services
- Provide self-diagnosing logs

This transforms the project into a fully autonomous resource manager.

7. Security and Anomaly Detection Enhancements

Future updates can incorporate:

- Threat detection (e.g., malicious CPU spikes)
- Behavior-based anomaly detection
- Process sandboxing

This would improve system reliability against suspicious processes.

8. Cloud and Edge Deployment

The system can be redesigned as:

- A scalable cloud service
- An edge device resource manager
- A portable API for IoT ecosystems

This would increase adoption across multiple computing platforms.

8. References:

- Tanenbaum, A. S. — *Modern Operating Systems*
- Linux Manual Pages: cgroups, nice, rlimit
- Python psutil documentation
- Prometheus and Grafana documentation
- Research papers on dynamic scheduling and RL optimizers

APPENDIX:

Appendix A – AI-Generated Project Elaboration/Breakdown Report

(This section contains all of the detailed explanation you received earlier, now compiled into one cohesive block for submission.)

1. Detailed Breakdown (Earlier AI Explanation)

Project Overview

The goal of this system is to dynamically monitor CPU & memory utilization in multiprogramming systems and adjust allocation to prevent bottlenecks. Expected outcomes include improved CPU utilization, reduced delays, and a real-time dashboard. Scope includes Linux user-space control but excludes kernel modification.

Module-Wise Breakdown

- Module A — Monitoring & Data Collection
Collect per-process data using psutil and expose metrics.
- Module B — Allocation Engine
Policy logic (rules or ML) + controller to modify cgroups/nice.
- Module C — Visualization
Dashboards using Grafana/Prometheus or custom frontend.

Functionalities

- Monitoring: CPU%, memory RSS, swap, I/O.
- Allocation Engine: heuristics, ML advisor, actuator safety.
- Dashboard: real-time charts, logs, manual override.

Technologies Recommended

- Python, psutil, Flask, scikit-learn, Prometheus/Grafana, Docker.
- cgroups v2 for process control.

Execution Plan

1. Build Monitoring Agent.
2. Implement rule-based controller.
3. Add visualization dashboard.
4. Add ML-based advisor (optional).
5. Testing using stress-ng.
6. Optimize hysteresis, cooldown, rollback.

Appendix B – Problem Statement

Adaptive Resource Allocation in Multiprogramming Systems

Description:

Develop a system that dynamically adjusts resource allocation among multiple programs to optimize CPU and memory utilization. The solution should monitor system performance and reallocate resources in real time to prevent bottlenecks.

Appendix C – Solution/Code

```
#include <bits/stdc++.h>

using namespace std;

// ----- PROCESS STRUCTURE -----
struct Process {
    int pid;
    int burstTime;
    int remainingTime;
    int priority;
    int memoryRequired;
```

```
int waitingTime = 0;
int turnaroundTime = 0;
string state; // Ready, Running, Waiting, Finished
};

// ----- SYSTEM METRICS -----
struct Metrics {
    int totalCPUTime = 0;
    int idleTime = 0;
    int totalMemory = 1024; // assume 1 GB memory
    int usedMemory = 0;
    float cpuUtilization = 0;
    float memoryUtilization = 0;
};

// ----- ADAPTIVE SCHEDULER -----
class AdaptiveScheduler {
private:
    vector<Process> processes;
    queue<int> readyQueue;
    Metrics metrics;
    int timeQuantum;
    int currentTime;

public:
    AdaptiveScheduler(int tq) {
        timeQuantum = tq;
        currentTime = 0;
    }

    void addProcess(Process p) {
```

```

processes.push_back(p);
readyQueue.push(p.pid);
metrics.usedMemory += p.memoryRequired;
}

void displayProcesses() {
    cout << "\nPID\tBurst\tRemain\tPriority\tMem\tState\n";
    for (auto &p : processes) {
        cout << p.pid << "\t" << p.burstTime << "\t" << p.remainingTime
            << "\t" << p.priority << "\t" << p.memoryRequired
            << "\t" << p.state << endl;
    }
}

void simulate() {
    cout << "\n==== Starting Simulation ====\n";

    while (!readyQueue.empty()) {
        int pid = readyQueue.front();
        readyQueue.pop();

        Process &p = processes[pid];
        p.state = "Running";
        cout << "\n[Time " << currentTime << "] Running Process " << p.pid << endl;

        int execTime = min(timeQuantum, p.remainingTime);
        currentTime += execTime;
        metrics.totalCPUTime += execTime;
        p.remainingTime -= execTime;

        // Update waiting times
    }
}

```

```

for (auto &q : processes) {
    if (q.state == "Ready") q.waitingTime += execTime;
}

// Adaptive logic: modify timeQuantum and priorities dynamically
adaptResources();

if (p.remainingTime == 0) {
    p.state = "Finished";
    p.turnaroundTime = currentTime;
    cout << "Process " << p.pid << " finished at time " << currentTime << endl;
} else {
    p.state = "Ready";
    readyQueue.push(p.pid);
}

updateMetrics();
displayProcesses();
}

showSummary();
}

void adaptResources() {
    // Adaptive logic: adjust time quantum or priorities
    int totalWaiting = 0, finished = 0;
    for (auto &p : processes) {
        if (p.state != "Finished") totalWaiting += p.waitingTime;
        else finished++;
    }
    int activeCount = processes.size() - finished;
}

```

```

int avgWaiting = (activeCount > 0) ? totalWaiting / activeCount : 0;

// Adjust time quantum dynamically
if (avgWaiting > 10) {
    timeQuantum = max(2, timeQuantum - 1);
} else if (avgWaiting < 5) {
    timeQuantum = min(8, timeQuantum + 1);
}

// Adjust priorities if someone is starving
for (auto &p : processes) {
    if (p.waitingTime > 15 && p.state == "Ready")
        p.priority++;
}

cout << "[Adaptive] New Time Quantum: " << timeQuantum << endl;
}

void updateMetrics() {
    metrics.cpuUtilization = (float)metrics.totalCPUTime /
        (float)(currentTime) * 100;
    metrics.memoryUtilization = (float)metrics.usedMemory /
        (float)metrics.totalMemory * 100;
}

void showSummary() {
    cout << "\n==== Simulation Summary ===\n";
    for (auto &p : processes) {
        cout << "Process " << p.pid
            << " -> Waiting: " << p.waitingTime
            << ", Turnaround: " << p.turnaroundTime << endl;
    }
}

```

```
}

cout << "\nCPU Utilization: " << metrics.cpuUtilization << "%";
cout << "\nMemory Utilization: " << metrics.memoryUtilization << "%";
cout << "\nSimulation Ended at time: " << currentTime << endl;
}

};

// ----- MAIN FUNCTION -----
int main() {
    AdaptiveScheduler scheduler(4);

    // Sample processes
    Process p1 = {0, 12, 12, 1, 200, 0, 0, "Ready"};
    Process p2 = {1, 8, 8, 2, 300, 0, 0, "Ready"};
    Process p3 = {2, 15, 15, 1, 250, 0, 0, "Ready"};

    scheduler.addProcess(p1);
    scheduler.addProcess(p2);
    scheduler.addProcess(p3);

    scheduler.displayProcesses();
    scheduler.simulate();

    return 0;
}
```